

AI GOVERNANCE — REGULATORY COMPLIANCE FRAMEWORK

Coaction Specialty Insurance *Based on NAIC Model Bulletin (December 2023) and NYDFS Proposed Circular Letter (January 2024) Applies to all business divisions, functions, and operational areas*

Our Commitment

At Coaction, we are accountable — to our regulators, our producers, and our policyholders — for every decision made with AI assistance anywhere in our business. That accountability does not diminish because a vendor built the model, because the process is operational rather than underwriting, or because the AI sits inside a system we bought rather than built. If Coaction's name is on the decision, Coaction owns the governance.

This framework applies to every AI system in use across Coaction — whether it prices a risk, processes a claim, reconciles cash, routes a document, validates a broker, or detects fraud. It covers predictive models, agentic automation platforms, neuro-symbolic systems, and any combination thereof. It covers systems we build, systems we buy, and systems operated on our behalf by third-party service partners.

The AI Landscape at Coaction

Coaction operates across three distinct AI categories, each with its own governance profile:

- 1. Predictive and Analytical Models** AI systems that analyze data to support underwriting, pricing, reserving, and fraud detection decisions. These are subject to the most direct regulatory scrutiny under NAIC and NYDFS guidance and require full actuarial defensibility.
- 2. Agentic AI (Operational Automation)** AI agents that autonomously execute multi-step operational tasks across any business function — including but not limited to claims processing, policy servicing, financial operations, document management, and producer services. These systems act on behalf of Coaction and are fully subject to this governance framework even where they do not directly influence underwriting or pricing.
- 3. Neuro-Symbolic Platforms** AI systems that combine neural pattern recognition with symbolic rule-based reasoning. These may operate across both analytical and operational domains simultaneously — for example, identifying litigation risk signals in claims while also routing correspondence. Both the neural and symbolic layers require separate, documented governance controls.

Any system that does not clearly fit one category should be treated as fitting all three until a determination is made.

What Needs to Be in Place

1. A Written AI Governance Program

Coaction must maintain a formal, board-approved written program that governs AI across the entire organization. This is not a policy document that lives in a shared drive — it is a live program with named owners, active processes, and documented evidence of operation.

The program must cover:

- Every AI system used across every function — underwriting, claims, finance, operations, producer services, compliance, and any other area where AI influences a process or outcome
- The full lifecycle of each system — from the decision to acquire or build, through deployment, ongoing monitoring, and retirement
- All three AI categories: predictive models, agentic platforms, and neuro-symbolic systems
- Both internally built systems and vendor-supplied or bought systems
- AI introduced into processes operated by third-party service partners on Coaction's behalf

Coaction may build this program on top of its existing Enterprise Risk Management framework and adopt the NIST AI Risk Management Framework as a structural foundation. We do not need to start from scratch — but we do need to extend existing governance infrastructure to genuinely cover our full AI footprint.

2. A Current AI Inventory

Coaction must maintain a complete, up-to-date inventory of every AI system in use, under active development, under evaluation, or recently retired — across every business function and division.

For each entry, the inventory must capture:

- What the system does and what processes or decisions it influences

- Which AI category it belongs to: predictive, agentic, neuro-symbolic, or combination
- What data it uses and where that data comes from
- Who owns it internally, and if externally supplied, who the vendor or partner is
- Its current deployment status and scope
- When it was last validated and when the next review is due
- Whether it operates within a third-party or BPO-managed process

This is the first document a regulator will ask for. If it is incomplete, inconsistent, or out of date, every other element of the governance program becomes harder to defend. Maintaining the inventory is an ongoing operational obligation, not a one-time exercise.

3. Documented Model and System Governance

For every AI system in the inventory, Coaction must maintain documentation that covers:

- The system's purpose, intended use, and the processes or decisions it is designed to support
- The data used to build or train it: source, lineage, quality, and suitability for the intended use
- How it was validated before deployment
- Its known limitations and the controls placed on its use
- How its performance is monitored on an ongoing basis
- What triggers a review, revalidation, or retirement decision
- How drift or degradation is detected and addressed

For predictive and analytical models, this includes a plain-language model narrative — a one-to-two-page document describing what the model does and how it works, written so that a senior leader or regulator with no technical background can read it and understand it.

For agentic AI systems, documentation must additionally cover:

- The defined scope of autonomous action — what the agent can and cannot do without human involvement

- The Human-in-the-Loop (HITL) design: what triggers human review, who reviews it, what authority they have, and how exceptions are tracked
- The exception rate threshold at which the system is flagged for review or remediation
- Change control procedures — how updates to agent logic, rules, or integrations are approved and documented

For neuro-symbolic systems, documentation must cover both layers independently:

- The neural layer: training data, validation approach, accuracy metrics, and drift monitoring
 - The symbolic layer: the documented rule set, the process by which rules are defined and updated, and who has authority to change them
 - The interaction between layers: how conflicts between neural outputs and symbolic rules are resolved, and how that resolution is audited
-

4. Human-in-the-Loop as Governance Control

For all agentic and neuro-symbolic AI systems operating across Coaction, Human-in-the-Loop (HITL) design is a primary governance control — not an afterthought.

Every agent deployment must document:

- The specific conditions under which human review are required before an action is taken or a decision is finalized
- The escalation path: who receives the exception, what information they are given, and what authority they hold
- The target exception rate and the threshold at which the system is reviewed
- How HITL performance is tracked over time and incorporated into revalidation

HITL design must be reviewed as part of every revalidation cycle. A system whose exception rate is trending toward zero requires as much scrutiny as one whose exception rate is too high — both indicate the system may not be operating as intended.

5. Ongoing Monitoring and Revalidation

Deploying an AI system is not the end of the governance obligation — it is the beginning. Coaction must maintain active monitoring for every AI system in production across all functions, including:

- Performance tracking against documented baselines established at deployment
- Detection of model or system drift — degradation in accuracy, reliability, or intended behavior over time
- A defined revalidation cadence, triggered by the passage of time, material changes to the system or data, or significant changes in the business or external environment
- A clear escalation path when monitoring flags a problem, with named accountability for resolution

For agentic systems, monitoring must also track operational outcomes — not just technical performance. If an agent is completing tasks accurately but producing outcomes that are inconsistent, inequitable, or contrary to Coaction's intent, that is a governance failure regardless of technical metrics.

6. Actuarial Validity of AI-Derived Variables

Any AI-derived factor used in underwriting, pricing, or reserving must be actuarially supportable. Coaction must be able to demonstrate:

- A clear, empirical, statistically significant relationship between the variable and the risk being assessed
- That the variable is not prohibited by applicable insurance law
- That its use produces consistent treatment of similarly situated risks

For commercial lines — and particularly for long-tail lines where AI-derived signals may influence reserve adequacy — this obligation is heightened. AI systems that influence reserving, including neuro-symbolic platforms used in claims severity assessment or litigation risk detection, are subject to the same actuarial defensibility standard as underwriting models.

7. Explainability Across All AI Categories

Coaction must be able to explain how each AI system operates in plain language that connects inputs to outputs in a logically intuitive way. This applies equally to:

- Predictive models used in underwriting and pricing
- Agentic systems routing documents, processing requests, or reconciling data
- Neuro-symbolic systems combining pattern recognition with rule-based logic

"The vendor's model produced this result" is not an acceptable answer — to a regulator, to an auditor, to a senior leader, or to a producer or policyholder who asks why a decision was made.

For neuro-symbolic systems specifically, explainability must extend to both layers: Coaction must be able to explain not just what the system decided, but which signals the neural layer identified and which rules the symbolic layer applied to reach that output.

8. Vendor and Partner Accountability

Coaction cannot outsource its AI governance obligations. For every third-party AI system, external data source, bought agent platform, or AI-enabled service operated by a partner on Coaction's behalf, the following must be in place:

- A documented due diligence process conducted before adoption, covering the vendor's model design, data practices, governance standards, and regulatory compliance posture
- Contractual terms establishing audit rights, requiring regulatory cooperation, mandating notification before material changes, and confirming Coaction's right to full documentation of how the system works
- Active exercise of those audit rights — not just having them on paper
- An obligation on the vendor or partner to operate to governance standards equivalent to those Coaction applies to its own systems

This obligation explicitly extends to:

- Platform vendors supplying agentic AI or neuro-symbolic systems
- BPO and operational service partners where AI is introduced into outsourced workflows — whether by Coaction or by the partner
- Data providers whose outputs feed AI models or agent decision logic

If a vendor's system produces an outcome that violates applicable law or Coaction's governance standards, that is Coaction's problem to remediate. Vendor relationships do not transfer accountability.

9. Board and Senior Management Accountability

AI governance at Coaction is an enterprise obligation owned at the top of the organization — not a technology function, not a compliance function, and not delegated to any single business unit.

Specifically:

- The board must approve AI governance policies and review them at least annually
- Senior management must establish and maintain the program, assign qualified owners across functions, oversee AI risk, and take prompt remedial action when problems are identified
- A cross-functional governance structure — spanning underwriting, claims, finance, operations, actuarial, data science, legal, compliance, and risk — must be in place and actively operating
- Regular, substantive reporting on AI risk and activity must reach the board

The governance structure must reflect the full scope of AI use across Coaction — not just the areas with the most obvious regulatory exposure. An operational AI agent processing 600 transactions per month carries governance obligations just as a pricing model does.

10. Training

All Coaction personnel who interact with, rely on, or are responsible for AI systems must be trained on the responsible and lawful use of those systems. Training must be:

- Tailored to each person's specific role and the AI systems they work with
- Completed promptly on onboarding and refreshed on a regular cadence
- Documented, with records maintained and available for examination

Training is not limited to technical staff. Senior leaders, underwriters, claims professionals, operations managers, and finance personnel who use AI-assisted outputs in their work are all within scope.

11. Examination Readiness

NAIC and NYDFS have signaled clearly that they will examine AI governance — and they have told us exactly what they will look for. Coaction must be able to produce on request:

- The written AI governance program
- The complete AI inventory
- System documentation for any AI system under review
- Evidence of validation, testing, and drift monitoring
- Vendor due diligence files and contracts
- Board minutes or reports evidencing governance oversight
- HITL design documentation for agentic systems
- Training records

The question is not whether Coaction will face an AI governance examination. It is whether we will be ready when it happens.

The Practical Test

Before representing Coaction's AI governance posture internally or to a regulator, every business and functional leader should be able to answer yes to these questions:

1. **Do we know what AI we are running?** Is the inventory complete, current, and maintained — across every function and division, not just underwriting?
2. **Can we explain what each system does and why?** Is there a plain-language narrative for every material AI system, covering all three AI categories?
3. **Do we know if our systems are still working as intended?** Is there active monitoring with documented drift detection, exception rate tracking, and escalation paths?
4. **Can we defend every AI-derived variable actuarially?** Are factors used in underwriting, pricing, and reserving supported by documented actuarial analysis?
5. **Are we genuinely accountable for our vendors and partners?** Do we have audit rights — and have we exercised them, including for agentic platform vendors and BPO partners?
6. **Are our HITL controls documented and functioning?** For every agent deployment, is the exception design documented, monitored, and reviewed?

7. **Could we hand a regulator our documentation today?** Is the governance program approved, evidenced, and examination-ready across the full scope of Coaction's AI use?

If the answer to any of these is no, that is a gap. Identifying the gap is the first step. Closing it — with documented, evidenced action — is the obligation.

This framework is maintained by AI Governance Lead at Coaction and reviewed annually by the Coaction management. Questions or governance concerns should be directed to [Contact].